RIO Release Notes

RIO 2.1.5

dsargent released this on Apr 10 · 18 commits to master since this release

Adds a feature that warns user when data is too noisy for RIO to improve predictions of the existing model. Also adds to the sample Notebook an example to demonstrate this warning.

Now the users can get more diagnosis information as a returned string after calling `train`.

RIO 2.1.4

dsargent released this on Mar 30 · 23 commits to master since this release

Hotfix for "path manipulation vulnerability" found by Fortify.

Will not impact user.

RIO 2.1.3

dsargent released this on Mar 4 · 25 commits to master since this release

Fixes installation via `pip`/PyPI/`egg`/`wheel` type mechanisms.

Will not impact user.

RIO 2.1.2

dsargent released this on Jan 29 · 49 commits to master since this release

Allow unlimited gRPC payload size.

Users have started encountering issues with large datasets exceeding the default gRPC payload size and getting errors.

This release allows for unlimited payload sizes. Although it has been tested with large and small datasets, it is possible that for some very large datasets other issues may start showing up, such as timeouts.

This removes the limitation on gRPC payload size. Now the users can send as much data as they want via the gRPC interface.

RIO 2.1.1

dsargent released this on Jan 8 · 54 commits to master since this release

Fixes Tensorflow "object not used" log spam.

This does not impact the user.

RIO 2.1.0

dsargent released this on Jan 7 · 57 commits to master since this release

Implements, and uses by default, ARD -- Automatic Relevance Determination (ARD), which means learning different lengthscales for each feature dimension.

This adds a new algorithmic feature, does not affect JPMC usages.

Rio 2.0.0

dsargent released this on Dec 10, 2019 · 74 commits to master since this release

This release is a stateless version of Rio. Models are no longer stored on the server side, and instead, are returned by the `train` API call. The client is expected to store the serialized model somewhere (in memory or on disk) and send the model to the server for subsequent `predict` calls.

The trained RIO model is now stored in the user side.

Rio "v2" -- separate training and prediction

dsargent released this on Nov 18, 2019 · 89 commits to master since this release

This is supposed to be a production-ready release of Rio. Major refactoring, in which the training step and prediction step are split into separate methods both in the gRPC interface and in the code itself.

This version involves gRPC interface when training RIO model and making predictions.

Checkpoint, including important bug fix

dsargent released this on Aug 13, 2019 · 101 commits to master since this release

Bug fix by Xin in training of neural net logic

This does not impact the end user.